

OS Programming

Files and File Systems

Some content for these slides comes from:
A Guide to Unix Using Linux, Fourth Edition

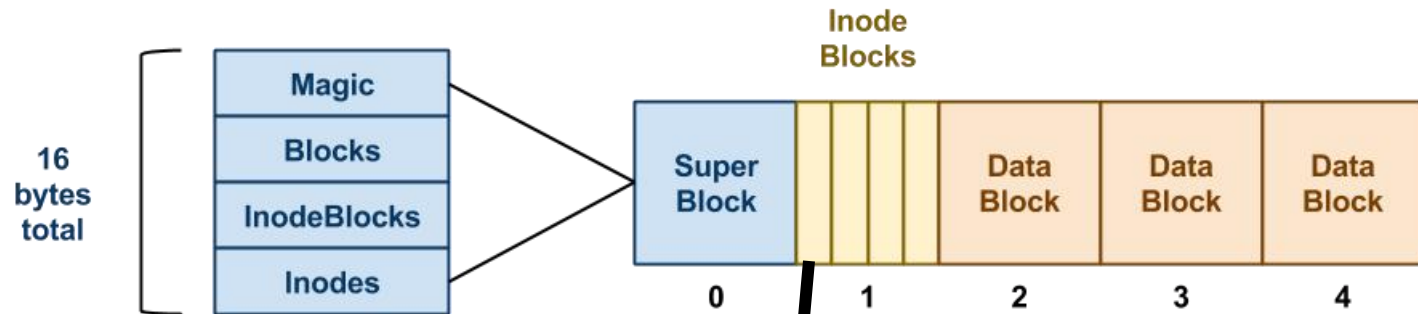
Understanding UNIX/Linux File Systems

- **File:** basic component for data storage
- **File system:** UNIX/Linux system's way of organizing files on storage devices
 - **Physical file system:** section of the hard disk that has been formatted to hold files
 - Disks are divided into **blocks**. A block is the smallest readable or writable unit which can be addressed. The size of these blocks varies.

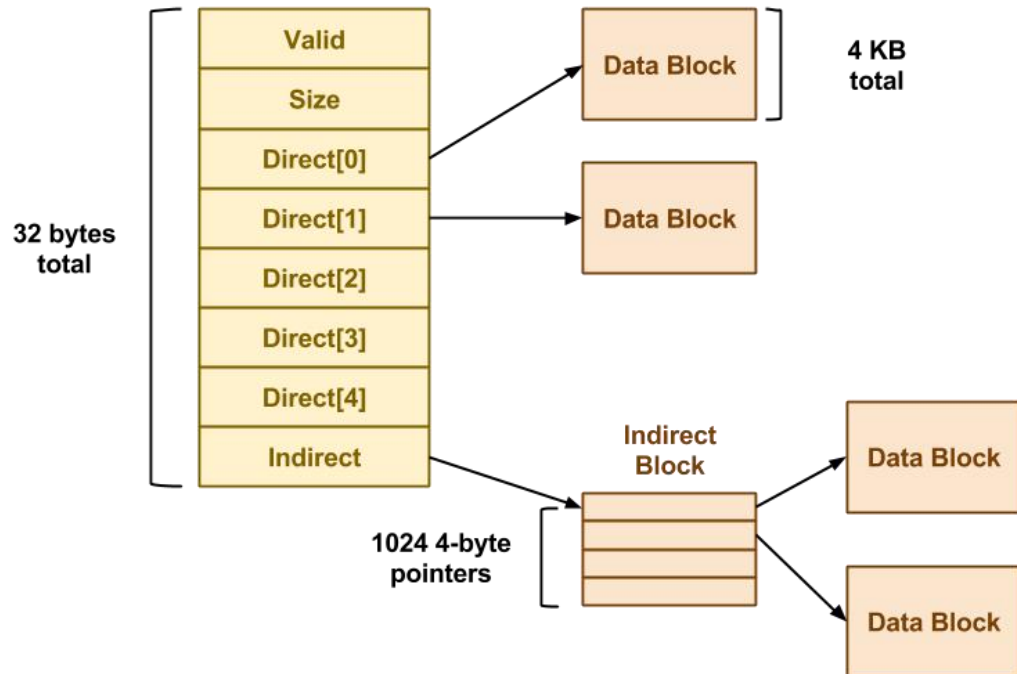
File System Concepts

- **Information nodes, or inodes**
 - Each directory/file has an inode and is identified by an inode number
 - Inode 2 contains the root of the directory structure (/)
 - Jumping-off point for all other inodes
 - inode 1 contains "bad blocks"
 - **ls -i** can be used to see the inode of files
 - Contains file/directory name, general information, pointer to the directory/file on a disk partition
- **Superblock** contains information about the layout of blocks on a specific partition

Basic Idea of a File System



Magic is a number that serves as a "signature" for the file system.



Understanding UNIX/Linux File Systems

- UNIX/Linux systems support many file systems
 - Examples: UNIX file system (ufs), extended file system (ext or ext fs)
- **ufs**: original native UNIX file system
 - Expandable, supports large amounts of storage, provides excellent security, reliable
 - Supports **journaling**
 - **The process of keeping chronological records of data or transactions so that if a system crashes without warning, the data or transactions can be reconstructed or backed up to avoid data loss or information that is not properly synchronized.**
 - Supports **hot fixes**
 - **The ability to automatically move data on damaged portions of disks to areas that are not damaged.**


Let's use **df -Th** to look at the filesystems for the online IDE.

Understanding UNIX/Linux File Systems


- In Linux, the native file system is **ext**
 - Installed by default
 - Modeled after ufs
 - First version contained some bugs
 - Newer versions of Linux use ext2, ext3, or ext4
 - ext4 enables the use of **extents**
 - An extent is a portion of a disk, such as a block or series of blocks, that is reserved for a file and that represents contiguous space, so that as the file grows, all of it remains in the same location on disk.
 - The use of extents reduces file fragmentation on a disk, which reduces disk wear and the time it takes to retrieve information.

Understanding UNIX/Linux File Systems

Extent-based vs. Block-based File Systems



Block-based: Meta-data stores a direct (pointer) to blocks of data.



Extent-based: Meta-data stores an extent (pointer and offset) to ranges of blocks of data.


Table 2-2 Comparison of typical file systems supported by UNIX/Linux

Feature	FAT	NTFS	ext4	ufs
Total volume or partition size	2 GB to 2 TB	2 TB	1 exabyte in Linux depending on the kernel version *	1 TB in Linux; 4 GB to 2 TB in UNIX depending on the version
Maximum file size	2 GB for FAT16; 4 GB for FAT32	Potentially 16 TB, but limited by the volume size (up to 2 TB)	16 GB to 2 TB in Linux depending on the kernel version *	2 GB in Linux; 2 GB to 16 TB in UNIX depending on the version
Security	Limited security based on attributes and shares	Extensive security through permissions, groups, and auditing options	Extensive security through permissions and groups	Extensive security through permissions and groups
Reliability through file activity tracking or journaling	None	Journaling	Journaling	Journaling
POSIX support	None (FAT16); limited (FAT32)	Yes	Yes	Yes
Reliability through hot fix capability	Limited	Supported	Supported	Supported
Support for extents	No	Yes, when pre-allocated via a program	Yes, when enabled	No
* These maximums are limited by the kernel version and are based on Linux kernel version 2.6.19.				

Understanding UNIX/Linux File Systems

- UNIX/Linux provides a layer of abstraction called **virtual storage**
 - Virtual storage can be allocated using different disks or file systems (or both), but that is transparently accessible as storage to users and programs.

FAT, EXT, and NFS are common filesystems.



Using UNIX/Linux Partitions

- **Partition:** section of disk that holds a file system
 - UNIX/Linux partitions identified with names
 - Examples: hda1, sda1
 - First two letters tell Linux the device type
 - Third letter indicates if disk is the primary or secondary disk
 - Partitions on a disk are numbered starting with 1

Setting Up Hard Disk Partitions

- Partition to organize space to contain file systems
- Some UNIX/Linux vendors recommend that:
 - Root partition holds the root file system directory
 - **Swap partition** acts like an extension of memory
 - General rule: same size as RAM
 - A swap partition enables **virtual memory**
 - **/boot partition** to store OS kernel files
- Mount partition to become part of file system

Using the mount Command

- Use *mount* to connect the file system partitions to the directory tree when the system starts

Syntax `mount [-option] [device-name mount-point]`

Dissection

- Use the *-t* option to specify a file system to mount.
 - *device-name* identifies the device to mount.
 - *mount-point* identifies the directory in which you want to mount the file system.
-

- **Example:**

```
mount -t iso9660 /dev/cdrom /media/cdrom
```

- Use *umount* before removing the storage media

```
umount /media/cdrom
```

Let's end with something hands-on

The tree, stat, and xxd commands

- **ls -R** and **tree** can be used to view a directory tree.
- The **stat** command shows information about a file, including the inode number.
- The **xxd** command can make a hexdump of a file or do the reverse.
 - a hexdump is a hexadecimal representation of the binary encoding of a file.

– example:

```
00000000 0000 0001 0001 1010 0010 0001 0004 0128
00000010 0000 0016 0000 0028 0000 0010 0000 0020
00000020 0000 0001 0004 0000 0000 0000 0000 0000
00000030 0000 0000 0000 0010 0000 0000 0000 0204
00000040 0004 8384 0084 c7c8 00c8 4748 0048 e8e9
00000050 00e9 6a69 0069 a8a9 00a9 2828 0028 fdfe
00000060 00fc 1819 0019 9898 0098 d9d8 00d8 5857
00000070 0057 7b7a 007a bab9 00b9 3a3c 003c 8888
00000080 8888 8888 8888 8888 288e be88 8888 8888
00000090 3b83 5788 8888 8888 7667 778e 8828 8888
000000a0 d61f 7abd 8818 8888 467c 585f 8814 8188
000000b0 8b06 e8f7 88aa 8388 8b3b 88f3 88bd e988
000000c0 8a18 880c e841 c988 b328 6871 688e 958b
000000d0 a948 5862 5884 7e81 3788 1ab4 5a84 3eec
000000e0 3d86 dcb8 5cbb 8888 8888 8888 8888 8888
000000f0 8888 8888 8888 8888 8888 8888 8888 0000
0000100 0000 0000 0000 0000 0000 0000 0000 0000
*
0000130 0000 0000 0000 0000 0000 0000 0000
000013e
```

Let's try them out!

Exercises

Create a file called "foo" that contains the word "hello".

Display the contents of the file "foo".

Get the inode number of the file "foo".

View a hexdump of "foo".